



Pipeline Operations

CS 4620 Lecture 10


Hidden surface elimination



Goal is to figure out which color to make the pixels based on what's in front of what.


Hidden surface elimination

- We have discussed how to map primitives to image space
 - projection and perspective are depth cues
 - occlusion is another very important cue




Back face culling

- For closed shapes you will never see the inside
 - therefore only draw surfaces that face the camera
 - implement by checking $\mathbf{n} \cdot \mathbf{v}$




Back-face Culling



- Problems ? 언제 문제가 생기나?
- Conservative algorithms
- Real job of visibility never solved



Painter's algorithm

- Simplest way to do hidden surfaces
- Draw from back to front, use overwriting in framebuffer




Painter's algorithm

- Amounts to a topological sort of the graph of occlusions
 - that is, an edge from A to B means A sometimes occludes B
 - any sort is valid
 - ABCDEF
 - BADCFE
 - if there are cycles there is no sort




Painter's algorithm

- examples of cycles
- No sort unless polygons are divided into small pieces



Painter's algorithm

- Useful when a valid order is easy to come by
- Compatible with alpha blending (later)






[Foley et al.]

The z buffer

- In many (most) applications maintaining a z sort is too expensive
 - changes all the time when the view changes
 - many data structures exist, but complex
- Solution: draw in any order, keep track of closest
 - Z-buffer keeps track of closest depth so far
 - when drawing, compare object's depth to current closest depth and discard if greater

Where Are We ?

- Canonical view volume (3D image space)
- Clipping done
- division by w
- $z > 0$




Z-Buffering: Algorithm

```
allocate z-buffer;           // Allocate depth buffer → Same size as
viewport.                    // viewport.

for each pixel (x,y)         // For each pixel in viewport.
    writePixel(x,y,backgrnd); // Initialize color.
    writeDepth(x,y,farPlane); // Initialize depth (z) buffer.


for each polygon             // Draw each polygon (in any order).
    for each pixel (x,y) in polygon // Rasterize polygon.
         $p_z$  = polygon's z-value at (x,y); // Interpolate z-value at (x, y).
        if ( $p_z > \text{z-buffer}(x,y)$ ) // If new depth is closer:
            writePixel(x,y,color); // Write new (polygon) color.
            writeDepth(x,y, $p_z$ ); // Write new depth.
```



Z-Buffering: Example

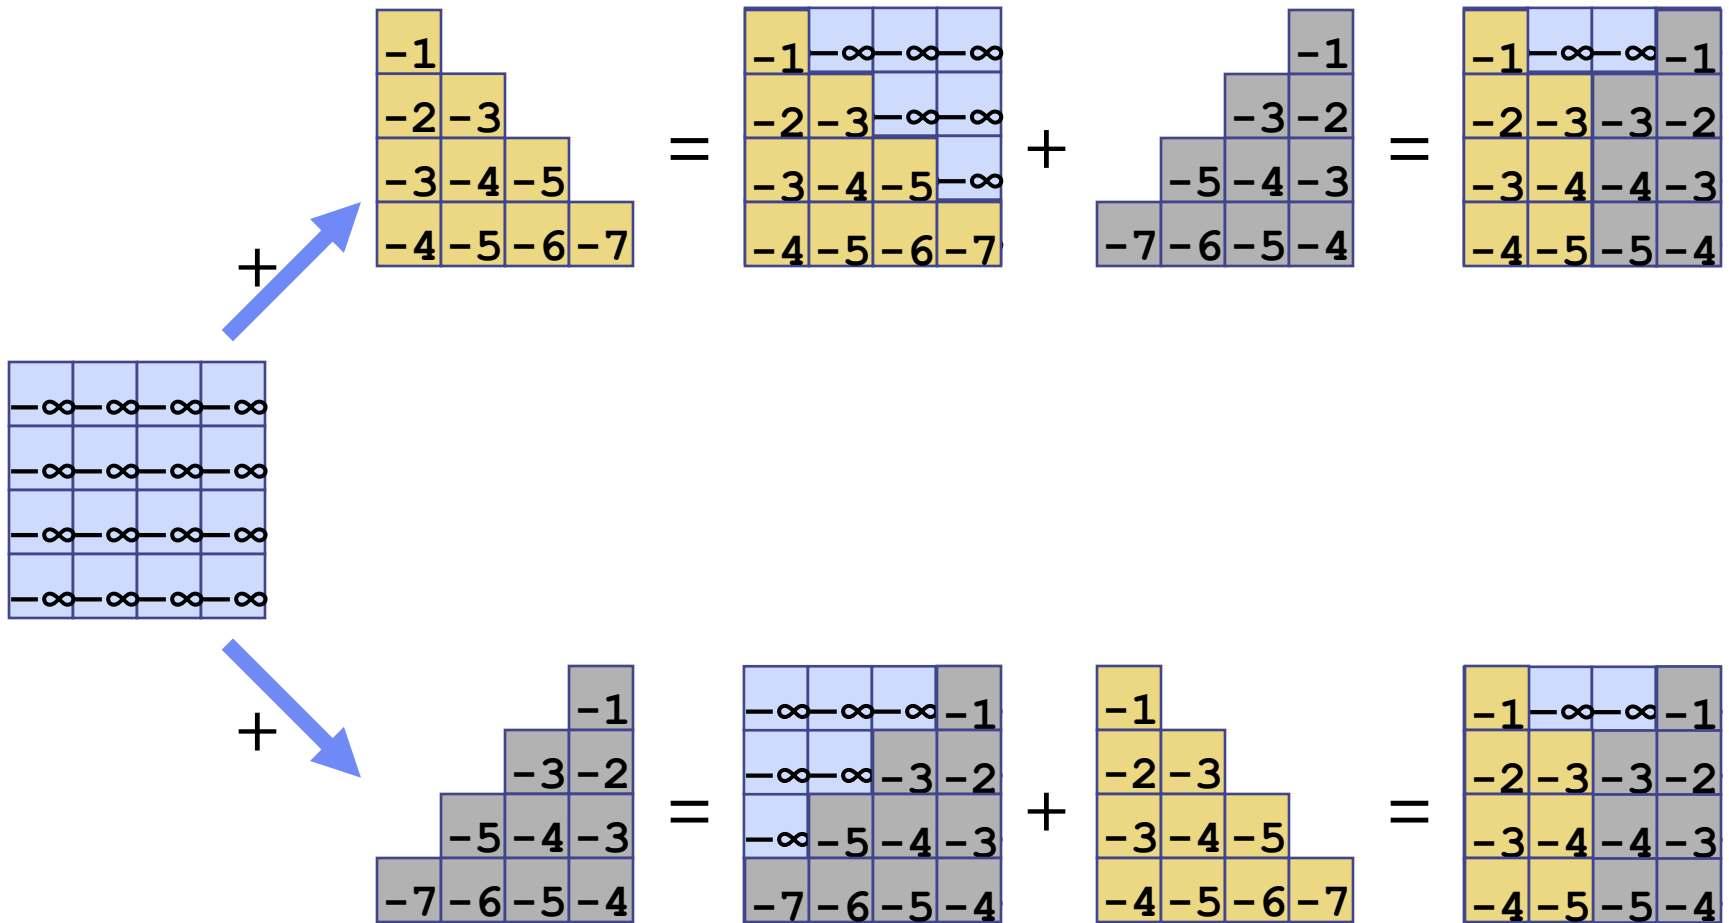
Scan convert the following two polygons.

The number in the pixel represents the z- value for that pixel




Does order matter?

Z-Buffering: Example



The z buffer



- another example of a memory-intensive brute force approach that works and has become the standard

Z-Buffering : Summary

- Advantages:
 - Easy to implement
 - Fast with hardware support → Fast depth buffer memory
 - On most hardware
 - No sorting of objects
 - Shadows are easy (later)
- Disadvantages:
 - Extra memory required for z-buffer:
 - Integer depth values
 - Scan-line algorithm
 - Prone to aliasing
 - Super-sampling


Pipeline for basic z buffer

- Vertex stage (input: position / vtx; color / tri)
 - transform position (object to screen space)
 - pass through color
- Rasterizer
 - interpolated parameter: z' (screen z)
 - pass through color
- Fragment stage (output: color, z')
 - write to color planes only if interpolated $z' <$ current z'

Precision in z buffer

- The precision is distributed between the near and far clipping planes
 - this is why these planes have to exist
 - also why you can't always just set them to very small and very large distances
- Generally use z' (not world z) in z buffer

Interpolating in projection



linear interp. in screen space \neq linear interp. in world (eye) space

Solution: linearly interpolate $1/z'$

Read, if you are interested, <http://www.lysator.liu.se/~mikaelk/doc/perspectivetexture/>

Flat shading

- Shade using the real normal of the triangle
 - same result as ray tracing a bunch of triangles
- Leads to constant shading and faceted appearance
 - truest view of the mesh geometry





Plate II.29 *Shutterbug*. Individually shaded polygons with diffuse reflection (Sections 14.4.2 and 16.2.3). (Copyright © 1990, Pixar. Rendered by Thomas Williams and H.B. Siegel using Pixar's PhotoRealistic RenderMan™ software.)

Pipeline for flat shading


- Vertex stage (input: position / vtx; color and normal / tri)
 - transform position and normal (object to eye space)
 - compute shaded color per triangle using normal
 - transform position (eye to screen space)
- Rasterizer
 - interpolated parameters: z' (screen z)
 - pass through color
- Fragment stage (output: color, z')
 - write to color planes only if interpolated $z' <$ current z'

Result of flat-shading pipeline



Local vs. infinite viewer, light

- Phong illumination requires geometric information:
 - light vector (function of position)
 - eye vector (function of position)
 - surface normal (from application)
- Light and eye vectors change
 - need to be computed (and normalized) for each face




Local vs. infinite viewer, light

- Look at case when eye or light is far away:
 - distant light source: nearly parallel illumination
 - distant eye point: nearly orthographic projection
 - in both cases, eye or light vector changes very little
- Optimization: approximate eye and/or light as infinitely far away


Directional light

- Directional (infinitely distant) light source
 - light vector always points in the same direction
 - often specified by position $[x \ y \ z \ 0]$
 - many pipelines are faster if you use directional lights




Infinite viewer


- Orthographic camera
 - projection direction is constant
- “Infinite viewer”
 - even with perspective, can approximate eye vector using the image plane normal
 - can produce weirdness for wide-angle views
 - Blinn-Phong: light, eye, half vectors all constant!



Specular shading (Blinn-Phong)



- Close to mirror \Leftrightarrow half vector near normal
 - Measure “near” by dot product of unit vectors



$$\mathbf{h} = \text{bisector}(\mathbf{v}, \mathbf{l})$$

$$= \frac{\mathbf{v} + \mathbf{l}}{\|\mathbf{v} + \mathbf{l}\|}$$

$$L_s = k_s I \max(0, \cos \alpha)^p$$

$$= k_s I \max(0, \mathbf{n} \cdot \mathbf{h})^p$$


↑
specularly
reflected
light

↑
specular
coefficient

Gouraud shading


- Often we're trying to draw smooth surfaces, so facets are an artifact

- compute colors at vertices using vertex normals
- interpolate colors across triangles
- “Gouraud shading”
- “Smooth shading”



[Gouraud thesis]


Plate II.30 *Shutterbug*. Gouraud shaded polygons with diffuse reflection (Sections 14.4.3 and 16.2.4). (Copyright © 1990, Pixar. Rendered by Thomas Williams and H.B. Siegel using Pixar's PhotoRealistic RenderMan™ software.)



Pipeline for Gouraud shading


- Vertex stage (input: position, color, and normal / vtx)
 - transform position and normal (object to eye space)
 - compute shaded color per vertex
 - transform position (eye to screen space)
- Rasterizer
 - interpolated parameters: z' (screen z); r, g, b color
- Fragment stage (output: color, z')
 - write to color planes only if interpolated $z' <$ current z'

Result of Gouraud shading pipeline



Vertex normals

- Need normals at vertices to compute Gouraud shading
- Best to get vtx. normals from the underlying geometry
 - e. g. spheres example
- Otherwise have to infer vtx. normals from triangles
 - simple scheme: average surrounding face normals



[Foley et al.]

$$N_v = \frac{\sum_i N_i}{\|\sum_i N_i\|}$$

Non-diffuse Gouraud shading

- Can apply Gouraud shading to any illumination model
 - it's just an interpolation method
- Results are not so models like specular
 - problems with any highlights smaller than a triangle

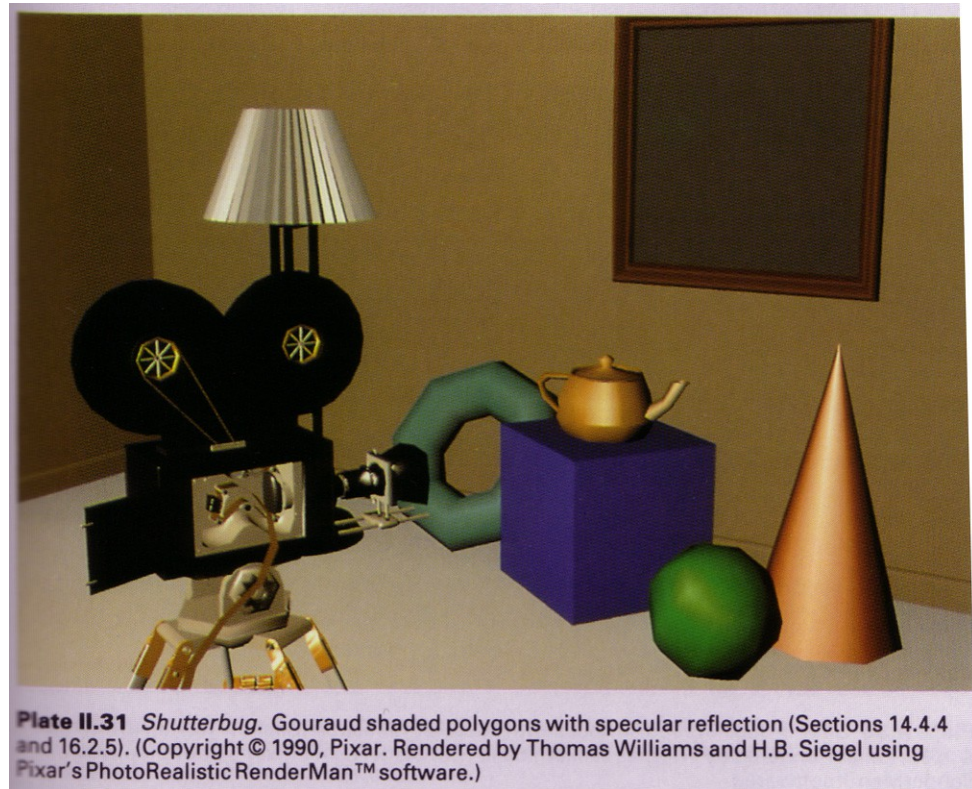



Plate II.31 *Shutterbug*. Gouraud shaded polygons with specular reflection (Sections 14.4.4 and 16.2.5). (Copyright © 1990, Pixar. Rendered by Thomas Williams and H.B. Siegel using Pixar's PhotoRealistic RenderMan™ software.)

Phong shading


- Get higher quality by interpolating the normal
 - just as easy as interpolating the color
 - but now we are evaluating the illumination model per pixel rather than per vertex (and normalizing the normal first)
 - in pipeline, this means we are moving illumination from the vertex processing stage to the fragment processing stage



[Foley et al.]

Phong shading

- Bottom line: produces much better highlights



Shutterbug. Gouraud shaded polygons with specular reflection (Sections 14.4.4 and 16.2.5). (Copyright © 1990, Pixar. Rendered by Thomas Williams and H.B. Siegel using Pixar's PhotoRealistic RenderMan™ software.)




Plate II.32 Shutterbug. Phong shaded polygons with specular reflection (Sections 14.4.4 and 16.2.5). (Copyright © 1990, Pixar. Rendered by Thomas Williams and H.B. Siegel using Pixar's PhotoRealistic RenderMan™ software.)

[Foley et al.]

Pipeline for Phong shading

- Vertex stage (input: position, color, and normal / vtx)
 - transform position and normal (object to eye space)
 - transform position (eye to screen space)
 - pass through color
- Rasterizer
 - interpolated parameters: z' (screen z); r, g, b color; x, y, z normal
- Fragment stage (output: color, z')
 - compute shading using interpolated color and normal
 - write to color planes only if interpolated $z' < \text{current } z'$

Result of Phong shading pipeline

